

FULL SCALE MINING

Jed Rembold

Wednesday, November 6, 2024

ANNOUNCEMENTS

- ⬡ Homework 8 due tomorrow
- ⬡ Test 2 results back on Monday
- ⬡ TechBytes tomorrow! Where does ethics belong in the data science curriculum?
- ⬡ Women in Tech: November 13th
 - ⬡ [RSVP here](#)
 - ⬡ Both speakers are awesome!
- ⬡ Polling today: polling.jedrebold.prof

REVIEW QUESTION

Which of the below phrases would match to the following regular expression:

```
[a-z]+-\d{2}$
```

- A) June-13-2022
- B) dec-8
- C) Jan-feb-28
- D) 9-11

REGULAR EXPRESSIONS IN SQL

BACK TO SQL

- ⬡ One of the main ways we previously used pattern matching was for filtering
- ⬡ You can also use regexes for pattern matching!
 - ⬡ `~` is a case sensitive match using the following regex
 - ⬡ `~*` is a case insensitive match using the following regex
 - ⬡ Either can have a `!` in front to negate the search (where things do **not** match the regex)

```
SELECT colname  
FROM tablename  
WHERE colname ~ '[a-z]*\s\d{2}';
```


EXTRACTING DATA

- ⬡ Another hugely common use of regex is to extract only the data you want from a much larger string
- ⬡ This can be particularly useful when cleaning data or constructing useful database tables
- ⬡ If you have just a single piece of information to extract, `SUBSTRING(text, regex)` is probably the most straightforward method
 - ⬡ If no capture group provided, the entire match is returned
 - ⬡ If a single capture group is provided, only that will be returned
 - ⬡ Just the **first** matching instance is returned

```
SELECT substring('today is November 6, 2024', '(\w+) \d{2}');
```

EXTRACTING MORE DATA

- ⬡ If you want to capture multiple groups at once, `regexp_match(str, regex)` is what you want
 - ⬡ What is returned is whatever is in any *capture groups* you may have included in your regex, or the entire match if there are no capture groups
 - ⬡ Output is returned as an array, to allow for potentially multiple capture groups
 - ⬡ You can index values out of the array or “unpack it” with `UNNEST`

```
SELECT regexp_match(
    'today is November 6, 2024',
    '(\w+) \d{2}.*(\d{4})')
);
```

REGULAR SPLITTING

- ⬡ You can also use regular expressions to replace or split text
- ⬡ `regexp_replace(text, re, replacement)` will replace the **first** matches of *re* in *text* with the replacement
 - ⬡ Add another `'g'` argument on the end to do so globally (replacing all matches)
- ⬡ `regexp_split_to_array(text, re)` will split the text into an array on each match to *re*
- ⬡ `regexp_split_to_table(text, re)` will split the text into a table column on each match to *re*. This is just like unnesting the array.

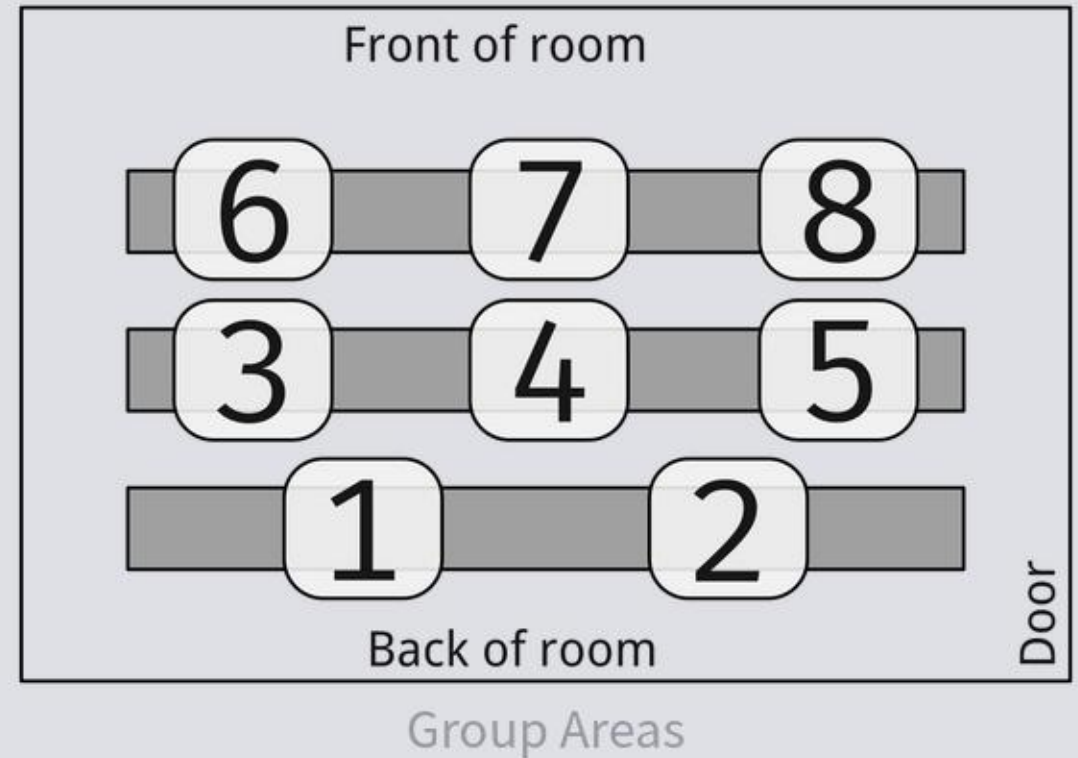


NOW TO YOU!

- ⬡ In the pairings below, take a look at [this](#) CSV file, which contains a simple subset of artists and dimensions from the MoMA data set
- ⬡ Create the simple table and import in the data from the CSV
- ⬡ See if you can achieve the following using regular expressions:
 - ⬡ Create and populate new columns for first, middle, and last name
 - ⬡ Create and populate new columns to hold the width and height in inches (in `1 1/4` form)
 - ⬡ If the above was easy and you finish quickly, create new columns to hold decimal equivalents of the width and height

TODAY'S GROUPS

- ⬡ Group 1: Dayton, Matthew, Haley
- ⬡ Group 2: Jordan, Marcus, Nick
- ⬡ Group 3: Tiffany, Mallory, Grace
- ⬡ Group 4: Harleen, Jack, Connor
- ⬡ Group 5: Aurora, Tippy, Sergio
- ⬡ Group 6: Jerrick, AJ, Michael
- ⬡ Group 7: Sam H, Greg, Evan
- ⬡ Group 8: Sam J, Hannah



FULL TEXT SEARCHING

FULL TEXT SEARCH

- ⬡ Pattern matching or regular expressions are great for parsing longer text for particular known patterns
- ⬡ What if you want more general information about the text within a longer passage?
 - ⬡ We need some different tools
- ⬡ Using these new tools will require using some new functions and data types:
 - ⬡ `tsvector`
 - ⬡ `tsquery`

TEXT TO VECTOR

- English (and other languages) have many connecting words that help convey meaning, intent, or relationships
 - to, on, at, with, of, etc.
- These are commonly not what are desired when searching text for particular ideas
- Instead it is useful to focus on *lexemes*
 - A lexeme is an abstract unit of meaning that underlies a set of words
 - RUN: run, runs, ran, running
- Postgres's `to_tsvector` will break a string down into its component lexemes, and keep track of where each occurred in the original string

```
SELECT to_tsvector('I flew back to Salem on Monday');  
>> 'back':3 'flew':2 'monday':7 'salem':5
```

LEXEME QUERIES

- ⬡ Text that you want to search **through** will need to be converted to a tsvector through `to_tsvector`
- ⬡ Text that you want to search **for** will be converted using `to_tsquery`
- ⬡ `to_tsquery` takes a sequence of words with symbols connecting them conveying relationships
 - ⬡ `&` - and
 - ⬡ `|` - or
 - ⬡ `!` - not
 - ⬡ `<->` - followed by

```
SELECT to_tsquery('fly & monday')
```

COMBINING VECTORS AND QUERIES

- ⬡ To actually complete a full text search, you ask Postgres to look through a `tsvector` object for a particular `tsquery`
- ⬡ The syntax to do so utilizes the match operator, which is two “at” symbols: `@@`

```
SELECT some tsvector @@ some tsquery;
```

- ⬡ Using `@@` is a true/false search, so the query is either found or not
 - ⬡ This means it can be used for filtering with `WHERE` as well!

INDEXING TSVECTORS

- ⬡ Individual `tsvector`s can not be easily ordered, so indexing a column with `tsvector` contents using the normal B-Tree method would not be effective
- ⬡ For this sort of content, use a *Generalized Inverted Index* or GIN indexing method instead

```
CREATE INDEX index name ON table  
USING GIN(column);
```

- ⬡ Seriously consider adding an index to your `tsvector` column, as it can *significantly* speed up these sorts of searches

GETTING MORE INFO

- ⬡ Sometimes it can be useful to get a bit more information about *where* a match shows up in the text
- ⬡ The `ts_headline` function can capture snippets of text around a match and display them
 - ⬡ `ts_headline` operates on the original text, **not** the tsvector! This will absolutely make it slower, so use it wisely!
- ⬡ There are a few required parameters and some options for `ts_headline`:
 - ⬡ the text to search
 - ⬡ the tsquery to look for
 - ⬡ Other options appear in an option string:
 - ⬡ `StartSel / StopSel`: the delimiters that will showcase the word
 - ⬡ `MinWords / MaxWords`: the min or max number of words to show around the match
 - ⬡ `MaxFragments`: a max number of fragments to show if the match occurs multiple times

RANKING

- ⬡ Sometimes a query might return many results, such that you would want a method to rank or compare them and only select the “best”
- ⬡ Ranking by relevancy is vague and tends to be very application specific, but it can be useful
- ⬡ `ts_rank` will give an arbitrary rank based on how many times your query words appear in the text
 - ⬡ This might mean that longer texts will always receive a higher rank! You can normalize by the text length by providing an extra numeric code as a third argument (2 or 8 probably best)
- ⬡ `ts_rank_cd` does similar, but also considers the proximity of searched lexemes
- ⬡ Both functions take 2 required arguments:
 - ⬡ a `tsvector` of the contents to be ranked
 - ⬡ a `tsquery` determining how they will be ranked

YOUR TURN!

The file [here](#) contains the SQL commands to generate and populate a simple table `alice` which hold the raw chapter contents of the book: Alice in Wonderland. You will need to set up your own `tsvector` column and index. In the same groups as earlier, see if you can use the data to answer the following:

- ⬡ In what chapters does the “Cheshire cat” appear?
- ⬡ In what chapter does the word “mushroom” appear the most? How many times does it appear?