

LOOKING THROUGH THE WINDOW

Jed Rembold

Monday, October 21, 2024

ANNOUNCEMENTS

- ⬡ HW7 due on Thursday night!
 - ⬡ Last of the “testable” content for the next exam
 - ⬡ The last part of the last question does involve some content from today
- ⬡ Exam 2 is 1 week from Wednesday
 - ⬡ Review and study materials will be posted on Wednesday
- ⬡ Polling today: polling.jedrembold.prof

REVIEW QUESTION!

Assuming that your server has its timezone set to 'American/Los_Angeles', what is the output of the below query?

```
SELECT make_time(  
    date_part('dow', current_date)::INT,  
    date_part('day', now())::INT,  
    date_part('hour', localtime)::INT  
) AT TIME ZONE 'PDT';
```

- A) 01:21:14-07
- B) 01:21:13
- C) 00:22:12-07
- D) 08:21:14

INTRO TO WINDOW FUNCTIONS

WINDOW FUNCTIONS

- ⬡ One useful piece of kit that is only partially included in the text (in Chapter 11) is that of *window functions*
- ⬡ A window function is like a mix between a normal column value and an aggregate function
 - ⬡ Unlike aggregate functions, a window function returns a value for each row
 - ⬡ Unlike normal column values, a window function can utilize other rows included within its “window” in making an aggregation
- ⬡ Any normal aggregate function can be used as a window function, though there are specific window functions as well
- ⬡ Window functions can only be used inside **SELECT** or **ORDER BY** statements
 - ⬡ They are evaluated *after* any filtering, grouping, or normal aggregations

OVER THE HILL

- ⬡ The defining characteristic of any window function is the **OVER ()** keyword, which comes after the aggregating window function
- ⬡ Content inside the **()** determines the “window” of the window function
- ⬡ By default, if nothing is provided, the entire column is the window
- ⬡ The below would output the average of the column in every row
 - ⬡ There will be other ways to do something similar with subqueries, but this is perhaps the most straightforward method

```
SELECT AVG(column) OVER()  
FROM table;
```

SPECIAL WINDOW FUNCTIONS

DEDICATED WINDOW FUNCTIONS

- ⬡ You can use any existing aggregate function as a window function, but there are also **more specific window functions**

Function	Description
<code>row_number()</code>	Assigns an ascending row number to each row in a window
<code>rank()</code>	Assigns an ascending rank to each row, with possible ties skipping the next value
<code>dense_rank()</code>	Assigns an ascending rank to each row, with possible ties not skipping the next value
<code>first_value(<i>col</i>)</code>	Returns the first value in the window of column <i>col</i>
<code>last_value(<i>col</i>)</code>	Returns the last value in the window of column <i>col</i>
<code>lag(<i>col</i>, <i>amt</i>)</code>	Returns the previous (or shifted by <i>amt</i>) row of column <i>col</i>
<code>lead(<i>col</i>, <i>amt</i>)</code>	Returns the next (or shifted by <i>amt</i>) row of column <i>col</i>
<code>nth_value(<i>col</i>, <i>n</i>)</code>	Returns the <i>n</i> th row of column <i>col</i> (<code>NULL</code> if doesn't exist)

CHANGING THE WINDOW

DETERMINING ORDER

- ⬡ Often, to be useful, you may want to define an ordering inside the `OVER()` statement
- ⬡ As soon as you specify an ordering, **the default window changes**
 - ⬡ By default, each window now encompasses everything from the start, up to (and including) that current row
 - ⬡ Easiest to see with classic aggregate functions

```
SELECT COUNT(*) OVER( ORDER BY column )  
FROM table;
```

TWEAKING THE WINDOW

- ⬡ You can tweak this window by specifying the starting and stopping point, using the syntax:

...*type* BETWEEN *offset* PRECEDING AND *offset* FOLLOWING

which appears in the **OVER** clause, after any provided ordering

- ⬡ *type* can be either **ROWS**, **RANGE**, or **GROUP**
- ⬡ *offset* can be
 - ⬡ an non-null, non-negative integer or **UNBOUNDED** if the type is **ROWS** or **GROUP**
 - ⬡ an value that makes sense to add or subtract from the ordered column if the type is **RANGE**

EXCLUDING THE WINDOW

- Can also exclude the current row or group from the window:

```
... EXCLUDE type
```

- type* here can be:
 - CURRENT ROW**, which excludes the current row from the window
 - GROUP** which excludes the current row and all other rows currently tied with it
 - TIES** which just excludes the tied rows, but keeps the current row

PARTITIONING

- ⬡ Additionally, you can specify a partition for the window
- ⬡ The default partition is the entire column
- ⬡ Partitioning here is determined *before* the rows comprising the window are computed
 - ⬡ This has a similar feel to `GROUP BY`, but *every* row will get a value here
- ⬡ Window functions evaluated within each window within each partition

```
SELECT AVG(column) OVER (  
    PARTITION BY column  
)  
FROM table;
```


ACTIVITY!

A SWEET DATASET

- ⬡ Contained [here](#) is an SQL file to generate a table of Halloween trick-or-treater data, which contains two columns:
 - ⬡ The timestamp when a trick-or-treater visited and was given a piece of candy
 - ⬡ What candy was given to the trick-or-treater
- ⬡ You started out with a supply of 50 of each type of candy
- ⬡ You will have several questions you are trying to answer.

TODAY'S GROUPS

- Group 1: Tippy, Dayton, Sam J
- Group 2: AJ, Myles, Matthew, Nick
- Group 3: Hannah, Grace, Michael
- Group 4: Aurora, Mallory, Greg
- Group 5: Jordan, Marcus, Evan
- Group 6: Harleen, Jerrick, Jack
- Group 7: Finn, Connor, Sam H
- Group 8: Haley, Tiffany, Sergio



QUESTIONS

- ⬡ Reminder: You started out with a supply of 50 of each type of candy
- ⬡ Using the dataset, answer the following questions:
 - ⬡ Between the 20th and 30th minutes, what were the three most popular candies given?
 - ⬡ What candies did you never run out of?
 - ⬡ At what time did you run out of Skittles?
 - ⬡ What times did you run out of the three most popular candies?
 - ⬡ Construct a 5-minute rolling average of the number of trick-or-treaters you saw each minute of the evening. According to this, about when were things busiest?